

# Why the Metropolis Algorithm Works

Here's a sketch of the important part of the **proof**, based on an argument in Gamerman (1997), of the **validity** of the Metropolis algorithm, in the case of a **discrete** (finite or countably infinite) state space  $S$  (see chapter 1 in Gilks et al. 1996 for a proof sketch when  $S$  is **continuous**).

I see now that my **Markov chain notation** up until this point has not been consistent enough to keep the proof from becoming confusing, so let's start over again with the following **notation**.

A **stochastic process**  $\{\theta_t^*, t \in T\}$ ,  $T = \{0, 1, \dots\}$  on a discrete state space  $S$  is a **Markov chain** iff

$$P(\theta_{t+1}^* = y | \theta_t^* = x, \theta_{t-1}^* = x_{t-1}, \dots, \theta_0^* = x_0) = P(\theta_{t+1}^* = y | \theta_t^* = x) \quad (30)$$

for all  $x_0, \dots, x_{t-1}, x, y \in S$ .

In general  $P(\theta_{t+1}^* = y | \theta_t^* = x)$  depends on  $x, y$ , and  $t$ , but if the probability of transitioning from  $x$  to  $y$  at time  $t$  is **constant** in  $t$  things will clearly be simpler; such chains are called **homogeneous** (confusingly, some sources call them **stationary**, but that terminology seems well worth avoiding).

The **random walk** described earlier is obviously a homogeneous Markov chain, and so are any Markov chains generated by the **MH algorithm**; I'll **assume homogeneity** in what follows.

Under **homogeneity** it makes sense to talk about the **transition probability**

$$P(x, y) = P(\theta_{t+1}^* = y | \theta_t^* = x) \quad \text{for all } t, \quad (31)$$

which **satisfies**

$$P(x, y) \geq 0 \text{ for all } x, y \in S \quad \text{and} \quad \sum_{y \in S} P(x, y) = 1 \text{ for all } x \in S. \quad (32)$$

# Metropolis Proof Sketch

When  $S$  is discrete a **transition matrix**  $P$  can be defined with element  $(i, j)$  given by  $P(x_i, x_j)$ , where  $x_i$  is the  $i$ th element in  $S$  according to whatever **numbering convention** you want to use (the second part of (32) implies that the row sums of such a matrix are always 1; this is the defining condition for a **stochastic matrix**).

Suppose the chain is **initialized** at time 0 by making a **draw** from a probability distribution  $\pi_0(x) = P(\theta_0^* = x)$  on  $S$  (**deterministically** starting it at some point  $x_0$  is a special case of this); then the probability distribution  $\pi_1(y)$  for where it will be at time 1 is

$$\begin{aligned}\pi_1(y) &= P(\theta_1^* = y) \\ &= \sum_{x \in S} P(\theta_0^* = x, \theta_1^* = y) \\ &= \sum_{x \in S} P(\theta_0^* = x) P(\theta_1^* = y | \theta_0^* = x) \\ &= \sum_{x \in S} \pi_0(x) P(x, y),\end{aligned}\tag{33}$$

which can be written in **vector** and **matrix** notation as

$$\pi_1 = \pi_0 P,\tag{34}$$

where  $\pi_0$  and  $\pi_1$  are regarded as **row vectors**.

Then by the **same reasoning**

$$\pi_2 = \pi_1 P = (\pi_0 P)P = \pi_0 P^2,\tag{35}$$

and in **general**

$$\pi_t = \pi_0 P^t.\tag{36}$$

For **simple** Markov chains this can be used to work out the **long-run** behavior of the chain as  $t \rightarrow \infty$ , but this becomes **algebraically prohibitive** as the **transition behavior** of the chain increases in **complexity**.

## Proof Sketch (continued)

In any case for **ergodic** Markov chains the limiting behavior  $\pi(y)$  is **independent** of  $\pi_0$  and is characterized by the relation

$$\pi(y) = \sum_{x \in S} \pi(x) P(x, y), \quad \text{or} \quad \pi = \pi P, \quad (37)$$

which defines the **stationary distribution**  $\pi$  of the chain.

As we've seen above, the hard bit in verifying the **validity** of the Metropolis algorithm is demonstrating that the Markov chain created by running the algorithm has the **correct stationary distribution**, namely the target posterior  $p(\theta|y)$ ; one way to do this is the following.

It's possible to imagine running any **homogeneous** Markov chain  $\{\theta_t^*, t = 0, 1, \dots\}$  with transition probabilities  $P(x, y)$  **backwards** in time.

This new **reverse-time** stochastic process can be shown also to be a **Markov chain**, although it may not be **homogeneous**.

If it **is** homogeneous, and if in addition the reverse-time process has the **same transition probabilities** as the original process, the Markov chain is said to be **reversible**; all such chains satisfy the **detailed balance equation**

$$\pi(x) P(x, y) = \pi(y) P(y, x) \text{ for all } x, y \in S. \quad (38)$$

It turns out that if there's a distribution  $\pi$  satisfying (38) for an **irreducible** Markov chain, then the chain is **positive recurrent** (and therefore **ergodic**) and **reversible**, and its **stationary distribution** is  $\pi$  (sum (38) over  $y$  to get (37)).

## Proof Sketch (continued)

In other words, if you're trying to create an **ergodic Markov chain** and you want it to have some **target stationary distribution**  $\pi$ , one way to achieve this goal is to ensure that the chain is **irreducible** and that its **transition probabilities**  $P(x, y)$  satisfy **detailed balance** with respect to the target  $\pi$ .

Any **reasonable** proposal distribution in the Metropolis algorithm will yield an **irreducible** Markov chain, so the interesting bit is to **verify detailed balance**; the argument proceeds as follows.

Consider a given **target distribution**  $p_x$  on  $S$ ; we're trying to construct a Markov chain with **stationary distribution**  $\pi$  such that  $\pi(x) = p_x$  for all  $x \in S$ .

The **Metropolis algorithm**—(15), with the special case of the **acceptance probabilities** (14) reducing to the **simpler form**  $\min\left[1, \frac{p(\theta^*|y)}{p(\theta_t|y)}\right]$  by the assumption of a **symmetric** proposal distribution—actually involves **two related Markov chains**: the (**less interesting**) chain that you could create by **accepting all proposed moves**, and the (**more interesting**) chain created by the **actual algorithm**.

Let  $Q(x, y)$  be any **irreducible** transition matrix on  $S$  such that  $Q(x, y) = Q(y, x)$  for all  $x, y \in S$ ; this is the transition matrix for the (**less interesting**) chain induced by the proposal distribution.

Define the (**more interesting**) chain  $\{\theta_t^*, t = 0, 1, \dots\}$  (the **actual** Metropolis chain) as having transitions from  $x$  to  $y$  proposed according to  $Q(x, y)$ , except that the proposed value for  $\theta_{t+1}^*$  is **accepted** with probability  $\min\left(1, \frac{p_y}{p_x}\right)$  and **rejected** otherwise, **leaving** the chain in state  $x$ .

## Proof Sketch (continued)

The **transition probabilities**  $P(x, y)$  for the **Metropolis chain** are as follows: for  $y \neq x$ , and denoting by  $A$  the event that the proposed move is **accepted**,

$$\begin{aligned} P(x, y) &= P(\theta_{t+1}^* = y | \theta_t^* = x) \\ &= P(\theta_{t+1}^* = y, A | \theta_t^* = x) + P(\theta_{t+1}^* = y, \text{not } A | \theta_t^* = x) \\ &= P(\theta_{t+1}^* = y | A, \theta_t^* = x) P(A | \theta_t^* = x) \\ &= P(\theta_{t+1}^* = y | A, \theta_t^* = x) P(A) \\ &= Q(x, y) \min\left(1, \frac{p_y}{p_x}\right). \end{aligned} \quad (39)$$

A **similar calculation** shows that for  $y = x$

$$P(x, x) = Q(x, x) + \sum_{y \neq x} Q(x, y) \left[1 - \min\left(1, \frac{p_y}{p_x}\right)\right], \quad (40)$$

but this is not needed to show **detailed balance** because (38) is **trivially satisfied** when  $y = x$ .

When  $y \neq x$  there are **two cases**:  $p_y \geq p_x > 0$  (I'll give details in **this case**) and  $0 < p_y < p_x$  (the other case follows **analogously**).

If  $p_y \geq p_x$ , **note** that  $\min\left(1, \frac{p_y}{p_x}\right) = 1$  and  $\min\left(1, \frac{p_x}{p_y}\right) p_y = \min\left(p_y, \frac{p_x}{p_y} p_y\right) = \min(p_y, p_x) = p_x$ ; **then**

$$\begin{aligned} p_x P(x, y) &= p_x Q(x, y) \min\left(1, \frac{p_y}{p_x}\right) = p_x Q(x, y) \\ &= p_x Q(y, x) = Q(y, x) \min\left(1, \frac{p_x}{p_y}\right) p_y \\ &= p_y P(y, x) \end{aligned} \quad (41)$$

and the proof of **detailed balance**, and with it the **validity** of the **Metropolis algorithm**, is **complete**.

# Directed Acyclic Graphs

BUGS achieves its **generality** by means of two ideas:


- (1) Viewing Bayesian models as **directed (acyclic) graphs** (DAGs).

The **conditional independence** nature of **Bayesian hierarchical models**—in which quantities in the model depend on things one layer higher in the hierarchy but no higher (e.g., in the NB10  $t$  model (23) the  $y_i$  depend on  $(\mu, \sigma^2, \lambda_i)$  but not on  $\nu$ )—lends itself to thinking of all quantities in such models as **nodes** in a **directed graph**.

A DAG can be thought of as a **picture** in which known and unknown quantities are represented either by **squares** (for knowns) or **circles** (for unknowns), connected by **arrows** (from the **parents** to the **children**) that indicate the direction of the stochastic dependence.

The **acyclic** assumption means that by following the directions of the arrows it's impossible to return to a node once you've left it, and **stacked sheets** indicate **repetition** (e.g., across conditionally IID data values).

Here's a DAG for the **NB10 model** based on the  $t$  distribution.



# Adaptive Rejection Sampling

(2) Employing **adaptive-rejection sampling** (Gilks and Wild, 1992) to generate the random draws from the full conditional distributions, when they don't have **simple recognizable forms**.

As we've seen, **rejection sampling** is a general method for sampling from a given density  $p(\theta|y)$ , which requires an **envelope function**  $G$  which dominates  $p$  (chosen so that  $G(\theta|y) \geq p(\theta|y)$  for all  $\theta$ ).

A restatement of the **algorithm** for normalized  $G$  (e.g., Ripley 1987) is

```
Repeat {  
  Sample a point theta from G ( . | y );  
  Sample a Uniform( 0, 1 ) random variable U;  
  If U <= p ( theta | y ) / G ( theta | y ) accept theta;  
}  
until one theta is accepted.
```

If  $p(\theta|y)$  is **expensive** to evaluate, time can be saved by identifying **squeezing functions**  $a(\theta|y)$  and  $b(\theta|y)$  with  $b(\theta|y) \leq p(\theta|y) \leq a(\theta|y)$ ; to use these, replace the **acceptance step** above (line 4 in the algorithm) by

```
If U > a( theta | y ) / G( theta | y ) reject theta;  
  else if U <= b( theta | y ) / G( theta | y ) accept theta;  
  else if U <= p( theta | y ) / G( theta | y ) accept theta.
```

**Adaptive rejection sampling** (ARS; Gilks and Wild 1992) is a relatively efficient method of **adaptive envelope construction** that works as a basis for Gibbs sampling if all of the full conditional densities are **log concave** (formally, a function  $p(\theta|y)$  of a vector argument  $\theta$  is log concave if the **determinant** of


$$\frac{d^2 \log g}{dy dy^T} \quad (42)$$

is **non-positive**).

## ARS (continued)

For univariate  $\theta$  the idea (see the figure on p. 14) is that an envelope function  $\log G_S(\theta|y)$  can be constructed on the log scale by drawing **tangents** to  $\log p(\theta|y)$  at each point in a given set of  $\theta$  values  $S$ .

An **envelope** between any two adjacent points is then constructed from the tangents at each end of the interval defined by the points:



the tangents form the envelope function on the log scale, & the secants form a lower squeezing function

The envelope is **linear on the log scale**, so rejection sampling on the original scale is performed with **scaled exponential** distributions (this can be done efficiently), and you get a **lower squeezing function** for free.

The useful thing about this idea is that the envelope can be constructed **adaptively**, by adding points to  $S$  as new  $\theta$  are sampled—thus the envelope **improves** as more samples are drawn.



## BUGS (continued)

BUGS uses a **hierarchy** of methods to sample from the full conditionals: it first tries to verify **conjugacy**; if that fails it then tries to verify **log concavity** of the full conditionals and uses **ARS** if so; and if that fails “classic” BUGS **quits** and winBUGS **switches over to (non-Gibbs) Metropolis-Hastings sampling**.

Log concavity includes many, **but not all**, distributions occurring in standard models, e.g., a uniform  $U(a,b)$  prior on the degrees of freedom parameter  $\nu$  in the NB10  $t$  model **fails** log-concavity.

In classic BUGS such distributions must be **discretized** (BUGS allows discrete variables to have 500 possible values, which generally leads to **quite accurate approximations**).

**Running classic BUGS.** You make **four kinds** of files:

- (1) a **program** file, with suffix `.bug`, containing the specification of your **model**;
- (2) one or more **data** files, with suffix `.dat`;
- (3) an **initial values** file, with suffix `.in`; and
- (4) a **command** file with suffix `.cmd`, containing instructions that specify the burn-in and monitoring phases.

Here's the **data file** in the NB10 example.

```
list( y = c(409., 400., 406., 399., 402., 406., 401., 403.,
  401., 403., 398., 403., 407., 402., 401., 399., 400., 401.,
    [ several lines omitted ]
  401., 407., 412., 375., 409., 406., 398., 406., 403., 404.),
grid = c(1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
  1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
    [ several lines omitted ]
  1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1 )
```

## BUGS (continued)

And here are the BUGS **program** (.bug) and **initial values** (.in) files in the NB10 example.

```
model nb10;

const

    n = 100, g = 100;

var

    mu, tau, u, grid[ g ], nu, y[ n ], sigma;

data in "nb10.dat";
inits in "nb10.in";

{

    mu ~ dnorm( 0.0, 1.0E-6 );
    tau ~ dgamma( 0.001, 0.001 ); # specifying the
    u ~ dcat( grid[ ] );          # prior distributions
    nu <- 2.0 + u / 10.0;

    for ( i in 1:n ) {

        y[ i ] ~ dt( mu, tau, nu ); # likelihood

    }

    sigma <- 1.0 / sqrt( tau );    # defining any other
                                   # quantities to be
                                   # monitored
}

Initial values

list( mu = 404.59, u = 30, tau = 0.04,
      seed = 90915314 )
```

# Implementation Details

Here are two BUGS **command** (.cmd) files in the NB10 example.

compile( "nb10-1.bug" )		compile( "nb10-1.bug" )
update( 1000 )		update( 2000 )
monitor( mu )		monitor( mu, 8 )
monitor( sigma )		monitor( sigma, 8 )
monitor( nu )		monitor( nu, 8 )
update( 5000 )		update( 40000 )
q( )		q( )

**Some Details.** (1) The priors: (a) I want to use a **diffuse** prior for  $\mu$ , since I don't know anything about the true weight of NB10 *a priori*.

The phrase  $\mu \sim \text{dnorm}(0.0, 1.0\text{E-}6)$  in BUGS-speak means that  $\mu$  has a Gaussian prior with mean 0 and **precision**  $10^{-6}$ , i.e.,  $\text{SD} = 1/\sqrt{\text{precision}} = 1,000$ , i.e., as far as I am concerned *a priori*  $\mu$  could be **just about anywhere** between  $-3,000$  and  $3,000$ .

(b) Similarly I want a **diffuse** prior for  $\sigma^2$ , or equivalently for the **precision**  $\tau = \frac{1}{\sigma^2}$ .

As we saw in the Poisson LOS case study, one popular **conventional** choice is  $\tau \sim \Gamma(\epsilon, \epsilon)$  for a small  $\epsilon$  like 0.001, which in BUGS-speak is said  $\text{tau} \sim \text{dgamma}(0.001, 0.001)$ .

This distribution is **very close to flat** over an extremely wide range of the interval  $(0, \infty)$ , although it does have a **nasty spike** at 0 (as  $\tau \downarrow 0, \Gamma(\epsilon, \epsilon)(\tau) \uparrow \infty$ ).

As noted earlier, the idea behind diffuse priors is to make them approximately **constant in the region in which the likelihood is appreciable**.

For this purpose it's useful to remember what the **frequentist answers** for  $\mu$  and  $\sigma$  would be, at least in the Gaussian model we looked at earlier.

Recall that the 95% confidence interval (CI) for  $\mu$  came out  $(403.3, 405.9)$ —so you can guess that the **likelihood** for  $\mu$  would be **non-negligible** in the range from (say) 402 to 407.

## Diffuse Priors

As for  $\sigma$  (or  $\sigma^2$  or  $\tau$ ), in the model  $(Y_i|\mu, \sigma^2) \stackrel{\text{IID}}{\sim} N(\mu, \sigma^2)$ , it's a standard result from **frequentist distribution theory** that in repeated sampling

$$\frac{(n-1)s^2}{\sigma^2} \sim \chi_{n-1}^2, \quad (43)$$

where  $s^2 = \frac{1}{n-1} \sum_{i=1}^n (y_i - \bar{y})^2$  is **random** and  $\sigma^2$  is **fixed**,  
from which

$$P_f \left[ A \leq \frac{(n-1)s^2}{\sigma^2} \leq B \right] = 0.99 \quad (44)$$

for  $A, B$  **such that**

$$P_f(\chi_{n-1}^2 \leq A) = P_f(\chi_{n-1}^2 \geq B) = 0.005. \quad (45)$$

Thus, using **Neyman's confidence trick**,

$$P_f \left[ \frac{(n-1)s^2}{B} \leq \sigma^2 \leq \frac{(n-1)s^2}{A} \right] = 0.99; \quad (46)$$

in other words,  $\left[ \frac{(n-1)s^2}{B}, \frac{(n-1)s^2}{A} \right]$  is a  
**99% confidence interval** for  $\sigma^2$ .

With the **NB10 data**  $n = 100$  and  $s^2 = 41.82$ , and you can use R to do this analysis:

```
> y
[1] 409 400 406 399 402 406 401 403 401 403 398 403 407 402 401 399 400 401
[19] 405 402 408 399 399 402 399 397 407 401 399 401 403 400 410 401 407 423
[37] 406 406 402 405 405 409 399 402 407 406 413 409 404 402 404 406 407 405
[55] 411 410 410 410 401 402 404 405 392 407 406 404 403 408 404 407 412 406
[73] 409 400 408 404 401 404 408 406 408 406 401 412 393 437 418 415 404 401
[91] 401 407 412 375 409 406 398 406 403 404
```


## More Details

```
> print( n <- length( y ) )  
[1] 100  
  
> print( s2 <- var( y ) )  
[1] 41.8201  
  
> qchisq( 0.005, 99 )  
[1] 66.5101  
  
> qchisq( 0.995, 99 )  
[1] 138.9868  
  
> ( n - 1 ) * s2 / qchisq( 0.995, 99 )  
[1] 29.78837  
  
> ( n - 1 ) * s2 / qchisq( 0.005, 99 )  
[1] 62.24904  
  
> qchisq( 0.005, 99 ) / ( ( n - 1 ) * s2 )  
[1] 0.01606451  
  
> qchisq( 0.995, 99 ) / ( ( n - 1 ) * s2 )  
[1] 0.03357015
```

So the conclusion is that the **likelihood** for  $\tau = \frac{1}{\sigma^2}$  should be **non-negligible** roughly in the region from about 0.015 to 0.035.

The figure below plots the prior distributions for  $\mu$  and  $\tau$  and **verifies their diffuseness** in the relevant regions.

## More Details



1. (c) As for the **prior** on  $\nu$ , you can tell from the normal qqplot of the NB10 data that the degrees of freedom parameter in the underlying  $t$  distribution is **fairly small**.

I'm going to use a **uniform**  $U(c_1, c_2)$  **prior**, where  $c_1$  is small but not too small (as noted earlier, with  $\nu < 2$  the variance is infinite, which is **not realistic** as a model for actual data) and  $c_2$  is big enough not to **truncate** the likelihood function (experience tells me that  $c_2 = 12$  will suffice; this can also be determined via MCMC **experimentation**).

Classic BUGS can't figure out how to sample from a continuous  $U(c_1, c_2)$  prior on  $\nu$ , however, so instead I've used a **discrete uniform prior** on a  $g = 100$ -point grid from 2.1 to 12.0 in steps of 0.1 (that's what `u ~ dcat( grid[ ] );`  
`nu <- 2.0 + u / 10.0;` does when `grid[ ]` is a vector of 100 1s).

## More Details

WinBUGS has a **more elegant solution** to this problem which we'll look at later.

(2) **Initial Values.** I can make fairly decent guesses at all the parameters as **starting values** for the Markov chain:

(a) The **sample mean** is 404.59, which should be close to the posterior mean for  $\mu$  in the  $t$  model;

(b) I'm just going to **guess** that  $\nu$  is around 5, which is specified by taking  $u = 30$ .

(c) Earlier I said that  $V[t_\nu(\mu, \sigma^2)] = \sigma^2 \left( \frac{\nu}{\nu-2} \right)$ , so with  $\nu \doteq 5$  and a **sample variance** of 41.82 you get  $\tau = \frac{1}{\sigma^2} \doteq 0.04$ .

**A Running Strategy.** With a problem like this with **relatively few parameters**, I often start off with a burn-in of 1,000 and a monitoring run of 5,000 and then look at the **MCMC diagnostics** (to be covered below).

The left-hand part of the table at the top of page 54 shows the BUGS **commands** that carry out this run.

You can either type in these commands **interactively** one at a time at the keyboard or put them in a `.cmd` file and run BUGS in the **background** (this is useful when you're interested in **simulating** the Bayesian analysis of many similar datasets for research purposes; the latest release of WinBUGS now also has this capability).

This run took about **5 minutes** on a not particularly fast workstation (a SunBlade 150 running Solaris Unix at 600 Mhz), which is actually fairly slow for a 3-parameter problem (the **discrete grid sampling** for  $\nu$  slows things down a lot).

# Classic BUGS Run

```
rosalind 61> bugs
```

```
Welcome to BUGS on 20 th Feb 2003 at 16:38:29
```

```
BUGS : Copyright (c) 1992 .. 1995 MRC Biostatistics Unit.
```

```
All rights reserved.
```

```
Version 0.603 for unix systems.
```

```
For general release: please see documentation for disclaimer.
```

```
The support of the Economic and Social Research Council (UK)  
is gratefully acknowledged.
```

```
Bugs>compile( "nb10-1.bug" )
```

```
model nb10;
```

```
[here BUGS just echoes the model shown on page 53]
```

```
}
```

```
Parsing model declarations.
```

```
Loading data value file(s).
```

```
Loading initial value file(s).
```

```
Parsing model specification.
```

```
Checking model graph for directed cycles.
```

```
Generating code.
```

```
Generating sampling distributions.
```

```
Checking model specification.
```

```
Choosing update methods.
```

```
compilation took 00:00:00
```

```
Bugs> update( 1000 )
```

```
time for 1000 updates was 00:00:47
```

```
Bugs>monitor( mu )
```

```
Bugs>monitor( sigma )
```

```
Bugs>monitor( nu )
```

```
Bugs>update( 5000 )
```

```
time for 5000 updates was 00:03:56
```

```
Bugs>q( ) # (output file created; more about this later)
```



# Practical MCMC monitoring and convergence diagnostics

Remember questions (3) and (4) awhile ago?—(3) **How large** should  $b$  and  $m$  be? (4) More generally, how do you know when the chain has reached **equilibrium**?

A **large body of research** has grown up just in the last eight years or so to answer these questions (some **good reviews** are available in Gelman et al. 1995, Gilks et al. 1995, and Cowles and Carlin 1996).

The theoretical bottom line is unpleasant: you **can't ever be sure you've reached equilibrium**, in the sense that every **MCMC diagnostic** invented so far has at least one example in which it failed to diagnose problems.

However, a collection of **four of the best diagnostics** has been brought together in a set of R functions called **CODA** by Best, Cowles, and Vines (1995) (downloadable from the R web site).

I will briefly discuss each of these in the context of the **NB10 analysis**.

**Geweke** (1992) proposed a simple diagnostic based on **time series** ideas.

Thinking of each column of the MCMC dataset as a **time series** (with iterations indexing time), he reasoned that, if the chain were in equilibrium, the **means** of the first (say) 10% and the last (say) 50% of the iterations should be **nearly equal**.

His diagnostic is a  $z$ -score for testing this equality, with a separate value for each quantity being monitored: Geweke  $z$ -scores **a lot bigger than 2 in absolute value** indicate that the mean level of the time series is **still drifting**, even after whatever burn-in you've already done.